

Chapter 6: Synchronization

Prof. Li-Pin Chang
National Chiao Tung University

Module 6: Synchronization

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Semaphores
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples
- Atomic Transactions

BACKGROUND.

Background

同步

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **count** that keeps track of the number of full buffers. Initially, count is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer. 用counter的方式來測buffer內容量

Producer

```
while (true) {  
    /* produce an item and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing buffer滿了  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

Consumer

```
while (1)
{
    while (count == 0)
        ; // do nothing buffer空了
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;
    /* consume the item in nextConsumed */
}
```

Race Condition

- `count++` could be implemented as

```
register1 = count  
register1 = register1 + 1  
count = register1
```

- `count--` could be implemented as

```
register2 = count  
register2 = register2 - 1  
count = register2
```

- Consider this execution interleaving with “count = 5” initially:
 - S0: producer execute `register1 = count` {register1 = 5}
 - S1: producer execute `register1 = register1 + 1` {register1 = 6}
 - S2: consumer execute `register2 = count` {register2 = 5}
 - S3: consumer execute `register2 = register2 - 1` {register2 = 4}
 - S4: producer execute `count = register1` {count = 6}
 - S5: consumer execute `count = register2` {count = 4}

Counter may be 4 or 6, depends on the sequence of S4 and S5
Counter can even be 5

Race Condition #2

- 2 threads, sharing a variable “safe” which is true initially

Thread 1:

```
if( safe == TRUE)
{
    safe = FALSE
    ... Do something...
}
```

Thread 2:

```
if( safe == TRUE)
{
    safe = FALSE
    ... Do something...
}
```


THE CRITICAL-SECTION PROBLEM

do {

entry section

critical section

exit section

remainder section

} while (TRUE);

Solution to Critical-Section Problem

- **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
- **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely 無限期、不定期 延遲
- **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted 同時執行的數量有限制，多的會被丟進queue

Hardware-based approaches to process synchronization

Synchronization Hardware

- Many systems provide hardware support for critical section code
- Interrupt disabling
 - Uniprocessor: good
 - Multiprocessor: does not work
- Test and set or swap
 - Uniprocessor: works, but wastes CPU cycles
 - Multiprocessor: works

Interrupt Disabling

- Uniprocessor
 - Source of preemption: timer, IO completion
 - Masking interrupts prevents the running process from being preempted
- Multiprocessor 攔截中斷只能對一顆CPU生效
 - Masking the interrupt of a CPU does not prevent racing processes on the other CPUs from entering a critical section
- Privilege instruction, cannot be used in user mode!

Atomic Instructions

- Modern machines provide special atomic hardware instructions
 - Atomic = non-interruptable
- Test memory word and set value
- Swap contents of two memory words
- Also called “spin locks”

TestAndndSet Instruction

假設這是Atomic

- Definition (instruction behavior):

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```


Solution using TestAndSet

這樣難道就不會有Race問題?

- Shared boolean variable *lock*, initialized to false.

Solution:

```
do {
```

```
    while ( TestAndSet (&lock ))  
        ;    /* do nothing
```

```
    //    critical section
```

```
    lock = FALSE;
```

```
    //    remainder section
```

```
} while ( TRUE);
```

如果Lock裡面False

→傳回False，Lock裡改True

如果Lock裡面True

→回傳True，程式卡在while

Swap Instruction

- Definition:

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

Solution using Swap

簡單來說就是爭搶Lock的False值

- Shared Boolean variable *lock* initialized to FALSE;
Each process has a local Boolean variable *key*.

Solution:

```
do {
```

```
    key = TRUE;
```

```
    while ( key == TRUE)
```

```
        Swap (&lock, &key );
```

如果key是True
就不停地與lock換，如果lock false了，換出
來的key就會變false

```
        //      critical section
```

```
        lock = FALSE;
```

把Lock False掉，讓別人也能用

```
        //      remainder section
```

```
    } while ( TRUE);
```

base on swap

Spin lock in Linux

```
lock:                                # The lock variable. 1 = locked, 0 = unlocked.
    dd    0

spin_lock:    放1進eax
    mov     eax, 1                    # Set the EAX register to 1.

loop:
    xchg     eax, [lock]              # Atomically swap the EAX register with
                                     # the lock variable.
                                     # This will always store 1 to the lock, leaving
                                     # previous value in the EAX register.

    test     eax, eax                 # Test EAX with itself. Among other things, this will
    做AND                                     # set the processor's Zero Flag if EAX is 0.
    如果是0→搶到Lock                         # If EAX is 0, then the lock was unlocked and
    如果是1→沒搶到Lock                       # we just locked it.
                                     # Otherwise, EAX is 1 and we didn't acquire the lock.

    jnz      loop                    # Jump back to the XCHG instruction if the Zero Flag is
    如果是1就Loop                                     # not set, the lock was locked, and we need to spin.

    ret                                # The lock has been acquired, return to the calling
                                     # function.

spin_unlock:
    mov     eax, 0                    # Set the EAX register to 0.

    xchg     eax, [lock]              # Atomically swap the EAX register with
                                     # the lock variable.

    ret                                # The lock has been released.
```

TAS and SWAP

- Applicable to both uniprocessor and multiprocessor systems
- Can be used in user mode (and of course kernel mode)
- Problems
 - Wasting CPU cycles in uniprocessor system busy wating很浪費CPU時間
 - Because the contention is stateless, process starvation is possible 純粹比誰快搶到Lock → 可能有process會餓死(搶不到)

- The TAS/SWAP approach guarantees which one(s) of the following?
 - Mutual exclusive
 - Progressive 只要有空就能進去
 - Bounded waiting 沒有

A bounded-waiting solution based on TAS/SWAP

```
do {
    waiting[i] = TRUE;
    key = TRUE;
    while (waiting[i] && key)
        key = TestAndSet(&lock);
    waiting[i] = FALSE;

    // critical section

    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;

    if (j == i)
        lock = FALSE;
    else
        waiting[j] = FALSE;

    // remainder section
}while (TRUE);
```

Waiting[i]=TRUE:
Pi wishes to enter the critical section

Key=TRUE:
there is at least 1 process waiting

To pick up the next process in
waiting[] if there are any waiting
processes

離開時把下一個waiting[i] == true
的改成waiting[i] == false
這樣下一個就能直接進來

- The selection of the next process to go in is a part of the critical section
- Once a process wishes to go in, it will be selected by waiting at most n-1 processes, as waiting[] is visited circularly

Summary

- Uniprocessor
 - Interrupt disabling
 - Only available in kernel space
 - Increasing interrupt latency
 - Spin lock (test and set, swap)
 - Correct, but wasting CPU cycles
- Multiprocessor
 - Interrupt disabling
 - Does not work if two involved processes run on different processors
 - Spin lock
 - Correct, wasting CPU cycle in a minor degree

Pure-software approach to process synchronization (Peterson's solution)

Peterson's Solution

- Two-process solution
- Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.
- The two processes share two variables:
 - int `turn`;
 - boolean `flag[2]`
- The variable *turn* indicates whose turn it is to enter the critical section.
- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process P_i is ready to get in!

Algorithm for Process P_i & P_j

太難啦~不講
僅限兩個processes

P_i

do {

 flag[i] = TRUE;

 turn = j;

 while (flag[j] && turn == j);

 CRITICAL SECTION

 flag[i] = FALSE;

 REMAINDER SECTION

} while (TRUE);

P_j

do {

 flag[j] = TRUE;

 turn = i;

 while (flag[i] && turn == i);

 CRITICAL SECTION

 flag[j] = FALSE;

 REMAINDER SECTION

} while (TRUE);

Algorithm for Process P_i

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while ( flag[j] && turn == j);
```

CRITICAL SECTION

```
    flag[i] = FALSE;
```

REMAINDER SECTION

```
} while (TRUE);
```

Proof of

- Mutual exclusion

- flag[i], flag[j] are both true

- Turn is either i or j

- Will “turn==i” become invalid after P_i enters CS?
 - Impossible because only P_i itself do the change (i.e., $turn \leftarrow j$)

- Progressive

- (!!) P_i will enter the critical section when P_j has no interest in entering the critical section (i.e., flag[j]=FALSE)

- Bounded-waiting

- Consider an example

- $P_i \leftarrow P_j$ P_i wins


- P_i completes and P_i arrives again

- $\rightarrow P_i$ always gives the chance away first

信號機

SEMAPHORES

-- a general approach



Semaphore

- Synchronization tool that **does not require busy waiting**
- Semaphore S – integer variable
- Two standard operations modify S: **wait()** and **signal()**
 - Originally called **P()** and **V()**
- Less complicated
- Can only be accessed via two indivisible (atomic) operations

Semaphore Implementation with no Busy waiting

- A semaphore is associated with a **waiting queue**
 - Append blocked processes to the waiting queue
- Two operations:
 - block – place the process invoking the operation on the appropriate waiting queue.
 - **Running** → **waiting**
 - Wakeup – remove one of processes in the waiting queue and place it in the ready queue.
 - **Waiting** → **ready**

Semaphore Implementation

- Implementation of wait:

```
wait (S){  
    S--;  
    if (S < 0) {  
        add this process to waiting queue  
        block(); }  
}
```

- Implementation of signal:

```
Signal (S){  
    S++;  
    if (S <= 0) { 表示之前有人被卡在裡面  
        remove a process P from the waiting queue  
        wakeup(P); }  
}
```


Semaphore Implementation

- Semaphores themselves are critical sections
 - Techniques such as interrupt disabling or test-and-set, are used to implement `signal()` and `wait()`
 - plus a waiting queue

Semaphore as General Synchronization Tool

- Counting semaphore – integer value can range over an unrestricted domain
 - Negative runtime values are legit
 - Negative initial values are not allowed in POSIX, however
- Binary semaphore – integer value can range only between 0 and 1; can be simpler to implement
 - Also known as mutex locks
- Can implement a counting semaphore S using a binary semaphore; and vice versa
- Can provide
 - Mutex: init value = 1
 - Sequencing or event: init value = 0
 - Capacity control: initial value=capacity

根據初始值有不同的效果

Typical Usages of a Counting Semaphore

- The purpose of a semaphore can typically be determined by the initial value of the semaphore
- *Mutex lock*: init value = 1
- *Sequencing or event*: init value = 0 要等別人把這個semaphore叫起來
- *Capacity control*: initial value=capacity

初始值其實象徵了多少人能在critical section活動

Mutual exclusion

Semaphore mutex=1

Pi

```
do {  
    waiting(mutex);  
  
    // critical section  
  
    signal(mutex);  
  
    // remainder section  
}while (TRUE);
```

Pj

```
do {  
    waiting(mutex);  
  
    // critical section  
  
    signal(mutex);  
  
    // remainder section  
}while (TRUE);
```

Sequencing or event Semaphore $\text{synch}=0$

Pi

```
 $S_1$ ;  
signal(synch);
```

Pj

```
wait(synch);  
 $S_2$ ;
```

Capacity control
Semaphore $\text{sem} = \text{capacity}$

```
Pi, Pj, Pk, ...  
{  
    ...  
  
    wait(sem) ;  
  
    ...  
  
    signal(sem) ;  
    ...  
}
```

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1

P_0	P_1
wait (S);	wait (Q);
wait (Q); Deadlock	wait (S);
.	.
.	.
.	.
signal (Q);	signal (S);
signal (S);	signal (Q);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.
 - Waking up blocked processes in a first-come first-served manner is a solution

Classical Problems of Synchronization

Example #1

- [有酒食，先生饌] 小明和老師中午一起吃飯，請使用semaphore 幫助小明，讓他禮讓老師先用餐。

```
S=0;
ming()
{
    wait(S);
    // eat
}

professor()
{
    // eat
    signal(S);
}
```

Example #2

- [蜘蛛人] 傑克和羅絲相約去看電影，兩人約在電影院門口見面，如果有人先到的話，要等另一人到了，才可以進入電影院。

```
R=0;J=0
jack()
{
    signal(J);
    wait(R);
    ...
    // see the movie
}
ross()
{
    signal(R);
    wait(J);
    ...
    // see the movie
}
```

Example #3

- [睡成一片] 資工系期中考快到了，總共有1000位同學想進入自習室。因為座位有限，所以只能50 個人同時進入。

S=50

```
student()  
{  
    wait(S);  
  
    // go studying  
  
    signal(S);  
}
```

Example #4

- A DMA controller supports four channels of data transfer

S=4;T=1;c[4]={F,F,F,F};

proc()

{

wait(S);

wait(T);

// pick one unused channel among c[0],c[1],c[2],c[3]

// setup DMA transfer

signal(T);

// start DMA

// wait for DMA completion

signal(S);

}

Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem
- Sleeping Barber Problem

Bounded-Buffer Problem

- N buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
 - To protect the buffer
 - There can be **many** producers and consumers!!
- Semaphore **full** initialized to the value 0
 - 0 items (for the consumer)
 - Block on no items
- Semaphore **empty** initialized to the value N.
 - N free slots (for the producer)
 - Block on no free slot

Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
  
    // produce an item  
  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
} while (true);
```

Producers produce items
Consumers “produce” free slots

What are the initial values of
empty and full?

What happens if mutex is placed
in the outer scope?

Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
do {  
    wait (full);  
    wait (mutex);  
  
    // remove an item from buffer  
  
    signal (mutex);  
    signal (empty);  
  
    // consume the removed item  
  
} while (true);
```


Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can both read and write.
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time.
- Shared Data
 - Data set
 - Semaphore **mutex** initialized to 1. (to protect “readcount”)
 - Semaphore **wrt** initialized to 1.
 - Integer **readcount** initialized to 0.

The first R-W problem

- No readers will wait until the writer locked the shared object
- Readers need not to synch with each other

Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do {  
    wait (wrt) ;  
  
    //  writing is performed  
  
    signal (wrt) ;  
} while (true)
```

Initially, wrt=1 mutex=1

Readers-Writers Problem (Cont.)

- The structure of a reader process


```
do {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readercount == 1) wait (wrt) ;  
    signal (mutex)  
  
    // reading is performed  
  
    wait (mutex) ;  
    readcount -- ;  
    if readcount == 0) signal (wrt) ;  
    signal (mutex) ;  
} while (true)
```

1. No readers will wait until the writer locked the shared object
 2. Readers need not to synch with each other
- Simply using one mutex for R/W violates the second condition
 - If the first reader is blocked by wrt, then other readers are blocked by mutex
 - Otherwise, the writer is blocked
 - The writer may starve?

Readers-Writers Problem (Cont.)

- Mutex
 - Protect the data set and the “readcount”
- Wrt
 - Mutex, ensure mutual exclusion among
 - A writer
 - A writer
 - ...
 - A group of readers

Dining-Philosophers Problem



- Shared data
 - Bowl of rice (data set)
 - Semaphore `chopstick` [5] initialized to 1

Dining-Philosophers Problem (Cont.)

- The structure of Philosopher i :

```
Do {  
    wait ( chopstick[i] );  
    wait ( chopstick[ (i + 1) % 5] );  
  
    // eat  
  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (true) ;
```

Potential deadlocks!

Dining-Philosophers Problem (Cont.)

- Possible ways to prevent deadlocks
 - One person get the left stick first, the rest get the right stick first
 - Allow up to 2 people having stick(s)
 - Picking up two sticks simultaneously

理髮師
Sleeping Barber Problem



Sleeping Barber Problem

- A barbershop consists of awaiting room with n chairs and a barber room with one barber chair. If there are no customers to be served, the barber goes to sleep. server等待event
- If a customer enters the barbershop and all chairs are occupied, then the customer leaves the shop.
- If chairs are available but the barber is busy, then the customer sits in one of the free chairs.
 - If the barber is asleep, the customer wakes up the barber.

Sleeping Barber Problem

- Semaphore Customers = 0;
 - Event: there are waiting customers
 - The barber waits on it if there is no customer
- Semaphore Barber = 0;
 - Event: barber is ready
 - The customer waits on it if the barber is busy
- Semaphore accessSeats = 1;
 - int NumberOfFreeSeats = N; //total number of seats

Costumer's process

```
while(1) {           先鎖住信號機
    wait(accessSeats) //mutex protect the number of available seats

    if ( NumberOfFreeSeats > 0 )
    {
        NumberOfFreeSeats--; //sitting down on a chair
        signal(Customers) ;   //notify the Barber
        signal (accessSeats); //release the lock
        wait(Barber);          //wait if the B is busy
        ....                  //here the C is having his hair cut
    }
    else
    {
        //there are no free seats
        signal (accessSeats); //release the lock on the seats
        ...                  //C leaves without a haircut
    }
}
} //while(1)
```

Barber process

```
while(1) {
```

```
    wait(Customers)    ;    //wait for C and sleep
```

```
    wait (accessSeats); //mutex protect the number of  
                        // available seats
```

```
    NumberOfFreeSeats++; //one chair gets free
```

```
    signal(Barber);      //Bring in a C for haircut 把一個等待的Customer叫起來
```

```
    signal (accessSeats); //release the mutex on the chairs
```

```
    .....              //here the B is cutting hair
```

```
}//while(1)
```

Mutexes and Monitors

Mutex locks

- “MUTually EXclusive” access
- Conceptually equivalent to semaphores with initial value = 1
- Only the locker of a mutex can unlock the mutex 信號機沒這種限制
- APIs 基本上跟信號機初始值為1的一樣
 - `pthread_mutex_lock()`
 - `pthread_mutex_unlock()`
 - ...

Semaphores vs. mutexes

- `pthread_mutex_XXXX()` only for threads
 - `pthread.h`
 - Functionally equivalent to semaphore with init value=1
 - Applicable to threads only
- `sem_XXX()` processes & threads 都可以用
 - `semaphore.h`
 - Applicable to threads and processes
 - A semaphore can be signaled by any process/thread
 - `sem_wait()`, `sem_post()`, ...\

Problems with Semaphores

- Incorrect use of semaphore operations:
 - `signal (s) signal(s)` signal 中間沒有wait
 - `wait (s) ... wait (s)`
 - Omitting of `wait (s)` or `signal (s)`
- Monitor provides a higher level abstraction of critical section to avoid these programming errors

Monitors


- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...

    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
}
```

Schematic view of a Monitor

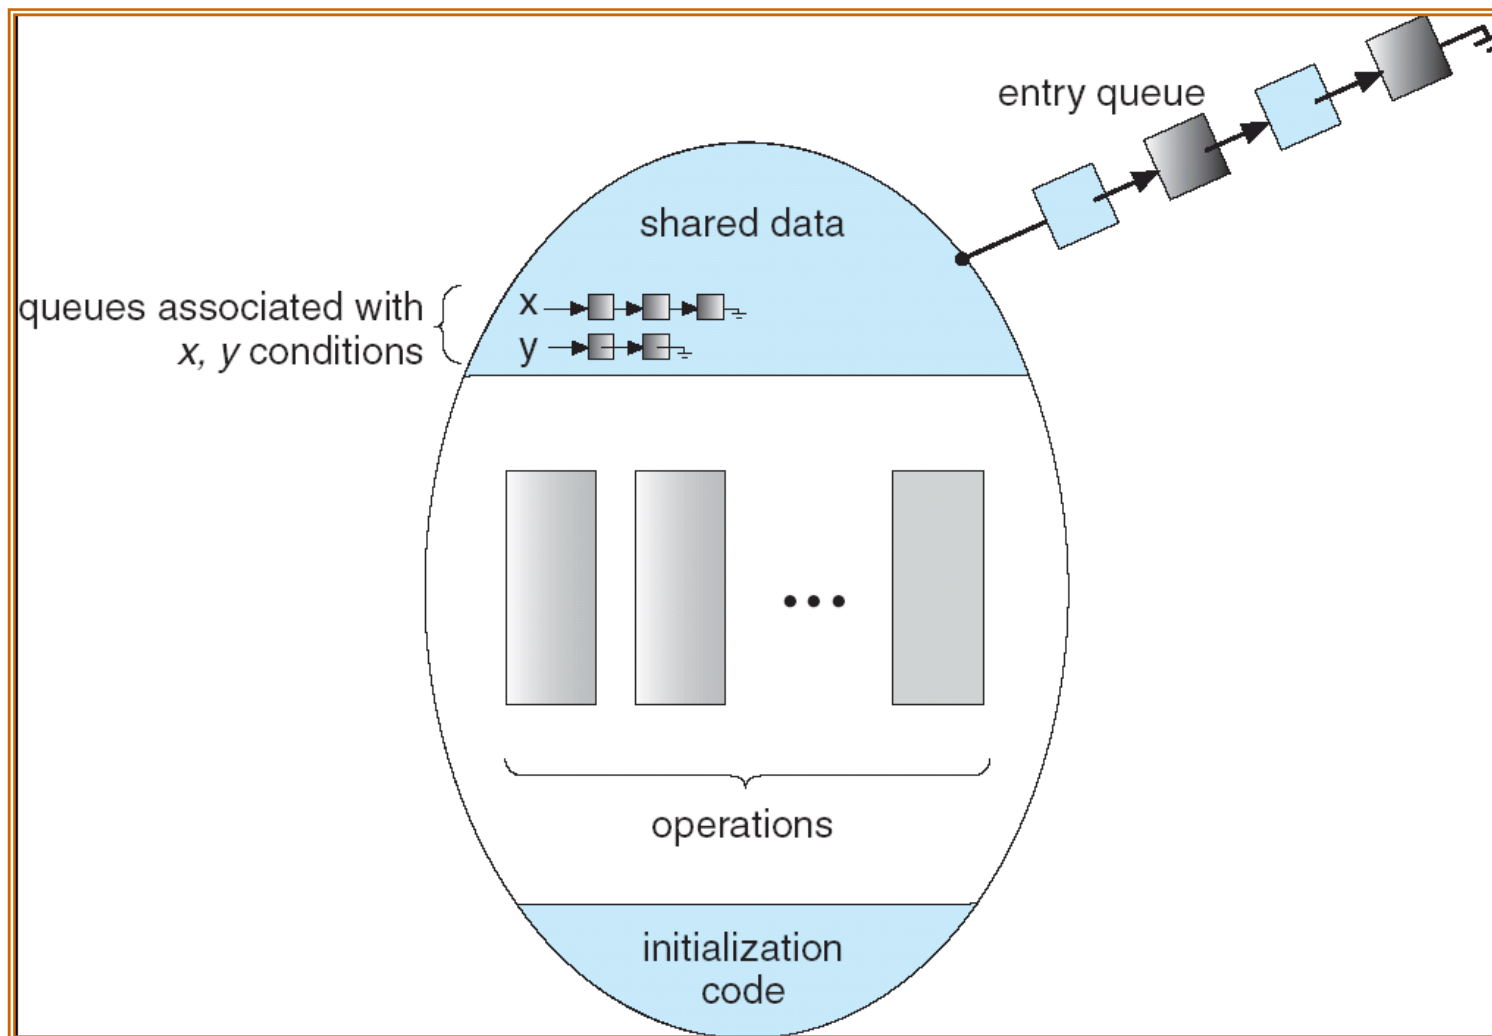


Condition Variables

- condition x, y;
- Two operations on a condition variable:
 - x.wait () – a process that invokes the operation is suspended. 掛在x上
 - x.signal () – resumes one of processes (if any) that invoked x.wait () 放出一個thread
- Monitor itself provides nothing but mutex
 - To implement other synch policy, conditional variables are needed
- signal → if there is no process waiting, nothing happens and the next process calls wait is blocked
 - Different from semaphore. For capacity control, a monitor must contain a counter

Monitor with Condition Variables

可以讓缺一些資料的thread等在variables上



Solution to Dining Philosophers

monitor DP

```
{  
    enum { THINKING; HUNGRY, EATING) state [5] ;  
    condition self [5];
```

```
void pickup (int i) {  
    state[i] = HUNGRY;  
    test(i);  
    if (state[i] != EATING) self [i].wait;  
}
```

```
void putdown (int i) {  
    state[i] = THINKING;  
    // test left and right neighbors  
    test((i + 4) % 5);  
    test((i + 1) % 5);  
}
```

Highlights to this solution:

- A philosopher picks up 2 chopsticks at a time
 - If he can not pick up 2 chopsticks, he waits
- After a philosopher done eating, he will check if his 2 neighbors can eat.

Solution to Dining Philosophers (cont)

```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}  
  
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}  
}
```

使用JAVA

```
BoundedBuffer(size, buffer) {  
    var buffer: Finite-Size Container;  
    var size, count: Integer;  
    var lock: Lock;  
    var notEmpty, notFull: Condition(lock); // cond. var. assoc. with "lock"  
  
    Produce(item) {  
        lock.acquire()  
  
        while (count == size) 如果buffer滿了就等  
            notFull.Wait()  
        buffer.add(item) // count++  
        notEmpty.Signal() 把consumer叫起來  
  
        lock.release()  
    }  
    Consume() {  
        lock.acquire()  
  
        while (count == 0)  
            notEmpty.Wait()  
        item = buffer.remove() // count--  
        notFull.Signal()  
  
        lock.Release()  
        return item  
    }  
}
```

The pairing acquire/release implements mutual exclusion of monitor

In Java, when you call a method with *synchronized* prefix, the acquire and release calls are automatically called (added to your code)

Java會自動幫你call lock

- How to implement semaphore using a monitor?
 - Signals to condition variables do not affect processes that are not in a monitor
 - Use a counter to keep track of “signals to semaphores”

SYNCHRONIZATION EXAMPLES

Synchronization Examples

- Solaris
- Windows XP
- Linux
- Pthreads

Solaris Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- Uses adaptive mutexes for efficiency when protecting data from short code segments
 - A process being blocked by a semaphore spins only if the process holding the semaphore is running on another CPU. Otherwise, the blocked process sleeps
- Uses condition variables and readers-writers locks when longer sections of code need access to data
- Uses turnstiles to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock

Windows XP Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems
- Uses spinlocks on multiprocessor systems
- Also provides dispatcher objects which may act as either mutexes and semaphores
- Dispatcher objects may also provide events
 - An event acts much like a condition variable

Linux Synchronization

- Linux:
 - disables interrupts to implement short critical sections
- Linux provides:
 - semaphores
 - spin locks

Pthreads Synchronization

- Pthreads API is OS-independent
- It provides:
 - mutex locks
 - condition variables
- Non-portable extensions include:
 - read-write locks
 - spin locks

ATOMIC TRANSACTIONS

Transaction Serializability

- A transaction is of a series of operations
- Even though operations to A and B are atomic, to T0 the operations to B may depend on the results of operations to A
- If transactions conflict, they must not be interleaved by one another

T_0	T_1
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)

T_0	T_1
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

Transaction Seriability

- Two-phase locking protocol
 - The growing phase → obtain all locks without releasing any locks
 - The shrinking phase → release all locks without obtain any locks
- Guarantees transaction seriability, but deadlocks are not avoided though

End of Chapter 6